

Matt Williams

mattwil230@gmail.com | San Francisco, CA | portfolio-eta-lilac-69.vercel.app

Full-Stack Engineer experienced across full-stack development, backend systems, APIs, and user-facing product interfaces, focused on shipping reliable, usable software around complex technical workflows.

Experienced across frontend development, data workflows, infrastructure, and deployment.

SKILLS

Frontend: React, Next.js, TypeScript, JavaScript, TailwindCSS

Backend: Python, FastAPI, Django, REST APIs, PostgreSQL, SQL, Supabase

Infrastructure: Docker, Linux, Git, GitHub, Vercel, RunPod

AI/ML: PyTorch, Hugging Face Transformers, LoRA/PEFT, CUDA Fundamentals, Model Evaluation

EXPERIENCE

Builda — Text-To-CAD Platform (March 2026 - Present)

- Built a full-stack text-to-CAD platform connecting a Fusion 360 UI add-in, FastAPI backend, REST API inference service, and GPU-backed model workflow to translate natural language into executable 3D CAD operations.
- Developed user-facing CAD generation workflows with structured output validation, schema checks, and algorithmic safeguards to convert model responses into safe Fusion 360 API actions.
- Fine-tuned and evaluated open-source LLMs using PyTorch, CUDA-enabled workflows, LoRA/PEFT, Unsloth, and Hugging Face Transformers, then integrated model outputs into a Dockerized FastAPI service with operational debug endpoints.

Crypto Asset Recovery — Machine Learning Engineer (February 2026 - Present)

- Built full-stack ML evaluation workflows using Python, Django/FastAPI, PostgreSQL, SQL, Pandas, NumPy, and PyTorch to review recovery candidates across a 130M+ row dataset, helping improve hash recovery success from 30% to 35%.
- Developed internal tools and backend data pipelines to preprocess 10M+ cleaned credential samples, inspect prediction results, and guide model iteration, improving top-N prediction accuracy by 12–18%.
- Built Dockerized evaluation and analysis tooling that allowed non-ML stakeholders to compare recovery candidates, review model outputs, evaluate failure patterns, and make product decisions from ML workflow results.

Unclone.io — Co-Founder (January 2026 - Present)

- Developed backend systems for an AI-powered clone-site detection platform, including signal ingestion, enrichment, similarity scoring, and verification workflows using Python, Django, Celery, PostgreSQL, and Docker.
- Designed model-assisted evaluation workflows to distinguish legitimate domains from malicious impersonation attempts, improving detection quality and reducing false positives.
- Built a human-in-the-loop remediation agent that generates provider-specific abuse reports and evidence bundles while enforcing human review before external submission.

Business Automation — AI Engineering Intern (November 2025 - February 2026)

- Developed QueryMango, a full-stack monitoring tool using Snowflake, SQL, Python, and alerting workflows to track 10+ database tables and detect pipeline failures in real time.
- Built data quality checks and pipeline health queries in Snowflake/SQL, identifying inconsistencies across 5–10% of datasets and improving visibility into failed data workflows.

STS Aerospace — Mechanical Engineering Intern (May 2024 - September 2024)

- Applied mechanical design, process optimization, and data analysis to improve manufacturing workflows, using SolidWorks, Excel, and Python to evaluate efficiency gains, cost savings, and production increases.

EDUCATION

University of New Hampshire, Durham, NH — B.S. in Mathematics, Minor in Mechanical Engineering